

Frozen reward comparison vs unguided: 4 independent batches $\times$ 50 candidates
Model affinity is not experimental binding; geometry evaluated separately

